# UNIT-4

# LEARNING

## LEARNING:

Learning in AI is a fascinating area! When it comes to AI, there are several different kinds of learning techniques, each with its own strengths and applications. Here's an overview of the main types of learning in AI:

### 1. Supervised Learning

- **Definition**: This type of learning involves training an AI model on a labeled dataset, meaning that for each input in the training set, the corresponding correct output is already known.
- **Example**: Training a model to recognize cats and dogs in images, where the dataset consists of images labeled as "cat" or "dog."
- **Common Algorithms**: Linear Regression, Decision Trees, Support Vector Machines (SVM), Neural Networks.

### 2. Unsupervised Learning

- **Definition**: Here, the model learns from data that has no labeled outputs. The goal is to find hidden patterns or structures in the data.
- **Example**: Customer segmentation in marketing, where an algorithm identifies groups of similar customers without pre-labeled categories.
- **Common Algorithms**: K-means Clustering, Hierarchical Clustering, Principal Component Analysis (PCA).

### 3. Reinforcement Learning

- **Definition**: In reinforcement learning, an agent learns to make decisions by interacting with an environment. The agent receives feedback in the form of rewards or penalties and aims to maximize its cumulative reward over time.
- **Example**: A robot learning to navigate through a maze by trial and error, or a game-playing AI like AlphaGo.
- **Common Algorithms**: Q-Learning, Deep Q-Networks (DQN), Proximal Policy Optimization (PPO).

### 4. Semi-Supervised Learning

- **Definition**: This approach combines a small amount of labeled data with a large amount of unlabeled data. It's used when labeling data is expensive or time-consuming.
- **Example**: A model trained with a few labeled medical images and many unlabeled ones to detect diseases in x-ray images.
- **Common Algorithms**: Label Propagation, Semi-Supervised Support Vector Machines (S3VM).

## Supervised Learning:

Supervised learning is a type of machine learning where the model is trained on **labeled data**. In this approach, the training dataset consists of input-output pairs, meaning that for each input (feature), there is a

known corresponding output (label). The model's task is to learn a function that maps inputs to the correct outputs based on the examples it is provided.

The goal is to **predict the output** for new, unseen inputs based on the patterns learned from the labeled dataset.

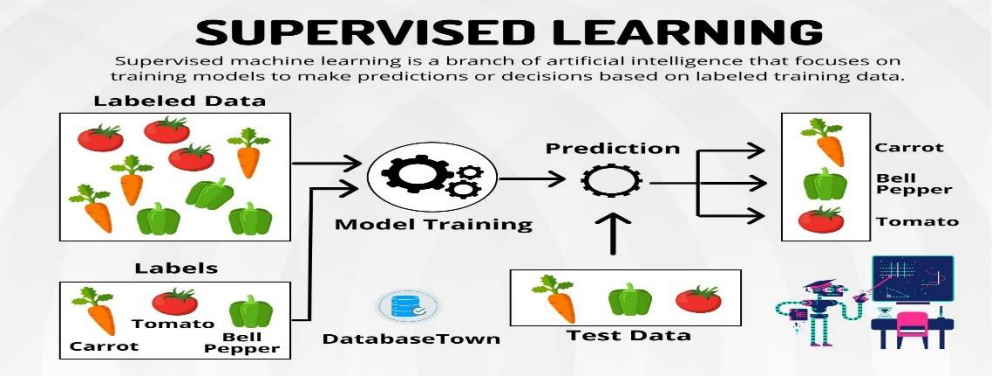


## How Does It Work?

1. **Training Phase**: The model is provided with a dataset where each example consists of an input and its corresponding output (label). The model learns the relationship between the input and output.
2. **Testing/Prediction Phase**: After training, the model is tested on new data (usually labeled as well). The model uses what it has learned to predict the output for this new data.
3. **Evaluation**: The model's performance is measured by comparing its predicted outputs to the actual outputs (labels) of the test dataset. Metrics like accuracy, precision, and recall are often used to assess the model's performance.

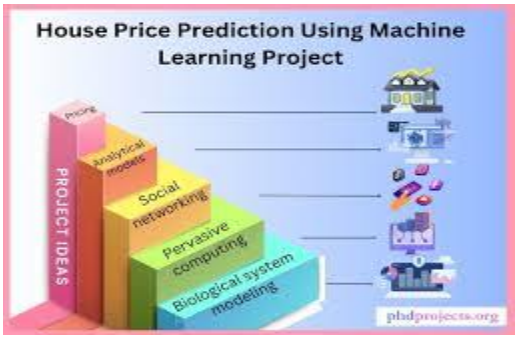


## Example: Predicting House Prices

Let's walk through an example where we use supervised learning to predict house prices based on certain features like size, location, and number of bedrooms.

**1. The Dataset (Training Data)**

Imagine we have a dataset that looks like this:

| House Size (sq ft) | Number of Bedrooms | Location | Price (in $1000) |
|---|---|---|---|
| 1500 | 3 | Suburb | 300 |
| 2000 | 4 | City Center | 500 |
| 1200 | 2 | Suburb | 250 |
| 1800 | 3 | City Center | 450 |

In this case:

- **Features (Inputs)**: House Size, Number of Bedrooms, Location

- **Label (Output)**: Price (the target we want to predict)

**2. The Goal**

The goal is to predict the **price** of a house (label) based on its **features** (size, number of bedrooms, and location). In supervised learning, we would feed the model this dataset and let it learn the relationship between the features and the price.

**3. Training the Model**

We train a supervised learning model (for example, a **linear regression** model) on this dataset. The model learns how house size, number of bedrooms, and location affect the price.

- For example, the model might learn that bigger houses tend to cost more and houses in the city center are more expensive than those in the suburbs.

**4. Making Predictions**

Once the model is trained, you can use it to predict the price of a new house. For instance:

- A house with **1600 sq ft**, **3 bedrooms**, and located in the **suburbs**.

The model will use the relationships it learned during training to estimate the price of this new house, such as predicting it might cost **$320,000**.

## Types of Supervised Learning Tasks

Supervised learning can be categorized into two types of problems:

1. **Regression**: Predicting a continuous output.
    - Example: Predicting the price of a house (as shown above), or predicting a person's weight based on their height and age.
2. **Classification**: Predicting a discrete label or category.
    - Example: Classifying emails as **spam** or **not spam**, or classifying images as **cats** or **dogs**.

### Advantages of Supervised Learning

- **Clear Objective**: Supervised learning has a clear target (the label) that the model is trying to predict.
- **Well-Defined Evaluation**: It's easy to evaluate the model's performance because you can compare the predicted outputs with the true labels.
- **Versatile**: It can be used for both regression (continuous outputs) and classification (discrete outputs) tasks.

### Disadvantages

- **Need for Labeled Data**: Supervised learning requires a large amount of labeled data, which can be expensive and time-consuming to obtain.
- **Overfitting**: The model might perform very well on the training data but fail to generalize to new data if it over fits (learns too much noise from the training set).

## Machine Learning:

**Machine Learning** is a branch of artificial intelligence (AI) that focuses on developing algorithms and models that enable computers to learn from and make predictions or decisions based on data, without being explicitly programmed. In machine learning, the system uses patterns or insights in the data to improve its performance on tasks over time.

### Key Points:

- It involves training a model using data.
- The model makes predictions or decisions based on patterns learned from that data.
- Machine learning can be **supervised**, **unsupervised**, or **reinforcement-based** depending on the problem and available data.

## Decision Trees:

A **Decision Tree** is a popular machine learning algorithm used for both **classification** and **regression** tasks. It's a type of **supervised learning** model that makes decisions based on the features of the data, and it does so in a tree-like structure, where each internal node represents a decision based on a feature, and each leaf node represents an outcome or prediction.

### How Decision Trees Work:

1. **Splitting**: The tree is built by recursively splitting the data into subsets based on the feature that best separates the data. The feature that maximizes some criterion (like information gain or Gini impurity) is chosen to split the data.
2. **Decision Nodes**: Each internal node in the tree represents a decision based on a particular feature. The data is split into branches depending on the feature values.
3. **Leaf Nodes**: These are the terminal nodes of the tree. Each leaf node corresponds to a predicted class (in classification) or a predicted value (in regression).
4. **Recursive Partitioning**: The tree-building process continues recursively, with the data being split further until the stopping condition is met (e.g., when the nodes are pure enough or a maximum depth is reached).

**Decision Tree Representation** is the structure of the decision tree model, which consists of **nodes** (decision points) and **edges** (branches that represent decisions). In a decision tree, the root node is the starting point, and the tree branches out into decision nodes and leaf nodes, which hold the predictions.

### Key Components of a Decision Tree:

1. **Root Node**: The very first decision point, where the data is split based on the best feature.
2. **Decision Nodes**: Intermediate nodes that further split the data based on another feature.
3. **Leaf Nodes**: Terminal nodes where the decision or prediction is made.
4. **Branches (Edges)**: Connections between nodes that represent the decisions or outcomes based on a specific feature.

### General Structure of a Decision Tree:

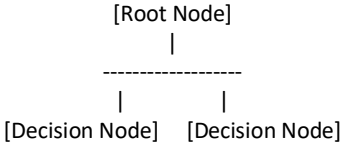

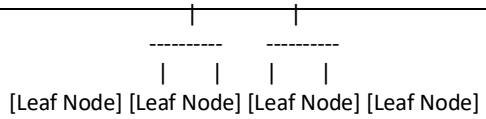


## Step-by-Step Representation of a Decision Tree:

1. **Root Node**: Start at the root. In the case of a classification problem, the root will ask a question about a feature in the dataset that provides the best split. For example, "Is the age greater than 30?"
2. **Decision Nodes**: After the root, we branch out based on the outcomes of the decisions made. For example, if the age is greater than 30, the tree will move to a different decision node asking a new question.
3. **Leaf Nodes**: The leaf nodes hold the outcome or decision. If it's a classification problem, the leaf node will indicate the class label. If it's a regression problem, the leaf node will show the predicted value.

## Example: Classifying Fruits Based on Weight and Color

Let's create a simple decision tree to classify fruits based on their **color** and **weight**. We have the following dataset:

| Fruit | Color | Weight (grams) |
|---|---|---|
| Apple | Red | 150 |
| Apple | Green | 120 |
| Orange | Orange | 180 |
| Orange | Orange | 160 |
| Apple | Red | 140 |

We want to **classify** the fruit as either "Apple" or "Orange" based on the color and weight.

## Step-by-Step Decision Tree Representation

1. **Start at the Root Node**:
   - The root node could be based on the **color** feature, since it's a strong distinguishing factor.
   - **Root Node**: "Is the color **Orange**?"
2. **Split the Data**:
   - If **Color = Orange**, the fruit is **Orange**.
   - If **Color ≠ Orange**, the fruit is an **Apple**.

This creates the following structure:

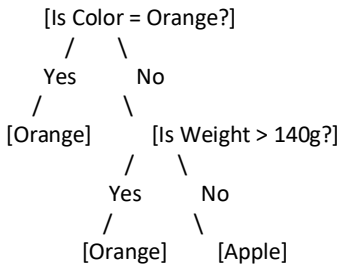


## Explanation of the Tree:

- **Root Node**: The tree first checks if the fruit's color is **Orange**.
  - If the **color is Orange**, the tree predicts the fruit is **Orange**.
  - If the **color is not Orange**, the tree moves to the next decision node.
- **Decision Node**: The second decision node checks the **weight** of the fruit.
  - If the **weight is greater than 140 grams**, the fruit is predicted as **Orange**.
  - If the **weight is less than or equal to 140 grams**, the fruit is predicted as **Apple**.

## INDUCING DECISION TREES FROM EXAMPLE

**Inducing Decision Trees from Examples** refers to the process of building a decision tree by using a set of training examples (data points) and recursively splitting them based on features, to maximize classification or regression accuracy.

Steps to Induce a Decision Tree

1. **Start with the Entire Dataset**: Begin with the whole dataset and create a decision tree based on the feature(s) that best split the data.
2. **Choose a Feature to Split**: For each node in the tree, choose the feature that best divides the data into subsets based on a specific criterion. Common criteria include:
   - **Gini Impurity** (for classification)
   - **Information Gain** (for classification)
   - **Mean Squared Error** (for regression)
3. **Split the Data**: Divide the dataset into subsets based on the chosen feature's values. Repeat the process recursively for each subset.
4. **Stop When**:
   - All data points in the subset belong to the same class (pure node).
   - There are no remaining features to split.
   - The tree reaches a predefined depth or size limit.
5. **Assign Predictions to Leaf Nodes**: Each leaf node in the tree is assigned the most frequent class (for classification) or the average value (for regression) of the data points that fall into that leaf.

## DECISION TREE LEARNING ALGORITHM

The **Decision Tree Learning Algorithm** is a supervised machine learning algorithm that builds a model (a decision tree) to predict the target variable by recursively partitioning the feature space. The goal of this algorithm is to divide the dataset into subsets that are as pure as possible with respect to the target variable, creating a tree-like structure of decision rules.

### Key Steps of the Decision Tree Learning Algorithm

The main steps of the decision tree learning process are:

1. **Select the Best Feature**: At each node of the tree, choose the feature that best splits the data based on a splitting criterion.

   Common criteria include:

   - **Information Gain (for Classification)**: Measures how much uncertainty is reduced in the data by splitting it on a particular feature.
   - **Gini Impurity (for Classification)**: Measures the "impurity" or disorder in the data after a split.

- **Mean Squared Error (for Regression)**: Used when predicting continuous values instead of classes.

2. **Split the Data**: Divide the dataset into subsets based on the selected feature. Each subset will correspond to a branch of the decision tree.
3. **Repeat Recursively**: For each subset, repeat the process of selecting the best feature and splitting the data. This recursive step continues until:
   - **Pure subsets**: The data at a node is homogeneous, meaning all examples belong to the same class (for classification) or have the same predicted value (for regression).
   - **Stopping Criterion**: Other stopping criteria might include reaching a maximum tree depth, a minimum number of samples per leaf, or a node with no further meaningful splits.
4. **Assign Class or Value to Leaf Nodes**: Once the data in a subset is pure or cannot be split further, assign a **class label** (for classification) or a **predicted value** (for regression) to the leaf node.

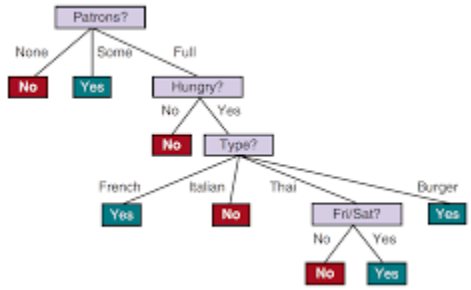


## REGRESSION AND CLASSIFICATION WITH LINEAR MODEL

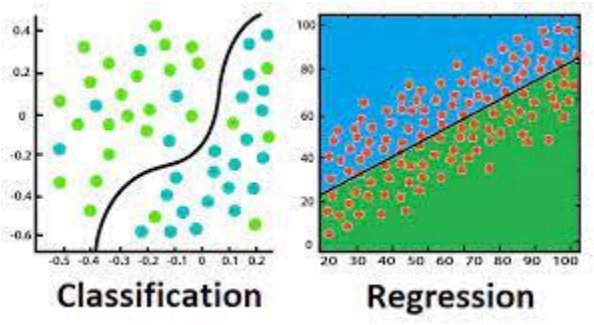


Regression using linear models is a method for predicting a **continuous numerical outcome** based on one or more input features. The core idea is to establish a **linear relationship** between the input variables (also called independent or explanatory variables) and the output variable (also called the dependent variable). The simplest form is **linear regression**, which fits a straight line to the data. The equation of this line is usually written as:

Y=w0+w1X1+w2X2+⋯+wnXn+ϵY = w_0 + w_1X_1 + w_2X_2 + \dots + w_nX_n + \epsilonY=w0+w1X1+w2X2+⋯+wnXn+ϵ

where YYY is the output variable, X1,X2,...,XnX_1, X_2, ..., X_nX1,X2,...,Xn are the features, w0w_0w0 is the intercept, w1,...,wnw_1, ..., w_nw1,...,wn are the weights (coefficients), and ϵ\epsilonϵ is the error term. The model tries to find the best weights that minimize the prediction error, often measured using **Mean Squared Error (MSE)**.

For example, consider a real estate company that wants to predict the **price of a house** based on features like its size in square feet, number of bedrooms, and the age of the house. By collecting data on past sales, a linear

regression model can be trained to learn the relationship between these features and the sale price. Once trained, the model can predict the price of new houses using the learned weights. If a house has 2000 sq ft, 4 bedrooms, and is 5 years old, the model will input these values into the equation and output an estimated price. This approach is widely used in industries like finance, marketing, and engineering, where understanding and forecasting continuous variables is critical.

## Classification with Linear Models:

Classification with linear models, such as **logistic regression**, is used when the goal is to predict **categorical outcomes**—for example, whether an email is spam or not, whether a patient has a disease, or whether a student will pass or fail. Unlike linear regression, which predicts real numbers, classification models predict **probabilities** that an input belongs to a particular class. In binary classification, logistic regression models the probability that the output belongs to class 1 using the **sigmoid function**, which maps any real-valued number into the range (0, 1):

P(Y=1|X)=11+e−(w0+w1X1+⋯+wnXn)P(Y=1|X) = \frac{1}{1 + e^{-(w_0 + w_1X_1 + \dots + w_nX_n)}}P(Y=1|X)=1+e−(w0+w1X1+⋯+wnXn)1

If the predicted probability is greater than 0.5, the instance is classified as class 1; otherwise, it's class 0. The model adjusts its weights to maximize the accuracy of classification based on a training dataset.

As an example, consider an email service provider that wants to detect **spam emails**. Emails are converted into numerical features such as the frequency of certain words ("free", "win", "money"), presence of links, or number of capital letters. These features are used to train a logistic regression model using a dataset labeled as “spam” or “not spam.” Once trained, the model can predict whether a new email is spam by computing the probability and comparing it to a threshold (usually 0.5). If the probability is high, the email is marked as spam and filtered out.

## Univariate Linear Regression:

**Univariate Linear Regression** is the simplest form of regression analysis that involves only **one independent variable** and one **dependent variable**. The purpose of this technique is to model the relationship between these two variables by fitting a straight line (called the **regression line**) through the data points. This line represents how the dependent variable (the one we want to predict) changes as the independent variable changes. The relationship is assumed to be **linear**, meaning as one variable increases or decreases, the other does so proportionally. The equation for this line is:

Y=w0+w1XY = w_0 + w_1XY=w0+w1X

where YYY is the predicted output, XXX is the input variable, w1w_1w1 is the slope (rate of change), and w0w_0w0 is the intercept (value of YYY when X=0X = 0X=0).

For example, suppose we want to predict the **price of a house** based only on its **size** (in square feet). By collecting data from past house sales, we may observe that larger houses tend to sell for higher prices. If we have data such as: a 1000 sq ft house sold for $150,000, a 1500 sq ft house sold for $200,000, and a 2000 sq ft house sold for $250,000, then univariate linear regression can find a straight line that best fits these points. Once the model learns the slope and intercept, it can predict the price of any house based on its size. For instance, for a house of 1800 sq ft, the model might predict a price of $230,000 based on the linear trend. This technique is widely used in fields like economics, real estate, and engineering for making simple but effective predictions.

## Multivariate Linear Regression:

**Multivariate Linear Regression** is an extension of simple (univariate) linear regression, where more than one **independent variable** is used to predict a **single dependent variable**. Instead of fitting a straight line to one feature, multivariate regression fits a **hyperplane** in a multidimensional space. This model is useful when the output depends on several factors working together. The general form of the equation is:

Y=w0+w1X1+w2X2+⋯+wnXn+ϵY = w_0 + w_1X_1 + w_2X_2 + \dots + w_nX_n + \epsilonY=w0+w1X1+w2X2 +⋯+wnXn+ϵ

Here, YYY is the predicted value, X1,X2,…,XnX_1, X_2, \dots, X_nX1,X2,…,Xn are the input variables (features), w0w_0w0 is the intercept, w1,w2,…,wnw_1, w_2, \dots, w_nw1,w2,…,wn are the coefficients (indicating the effect of each variable on the prediction), and ϵ\epsilonϵ is the error term.

For example, imagine you're working for a real estate agency and want to predict the **price of a house** not just based on its **size**, but also its **number of bedrooms** and **age**. A house's price might increase with size and number of bedrooms, but decrease as the house gets older. Suppose you have data showing that a 2000 sq ft, 3-bedroom, 5-year-old house sells for $350,000, while a 1500 sq ft, 2-bedroom, 15-year-old house sells for $220,000. Using this data, a multivariate linear regression model can learn how each factor affects the price. Once trained, the model can predict the price of any new house by plugging in its size, number of bedrooms, and age into the regression equation. This approach is highly valuable in real-world applications where multiple factors influence outcomes—such as predicting sales, estimating production costs, or forecasting performance in various domains.

## Linear classifiers with a hard threshold:

**Linear classifiers with a hard threshold** are a simple type of classification model that separates data into classes using a straight line (in 2D) or a hyperplane (in higher dimensions). The core idea is to compute a weighted sum of the input features and then apply a **hard threshold function** to decide the class label.

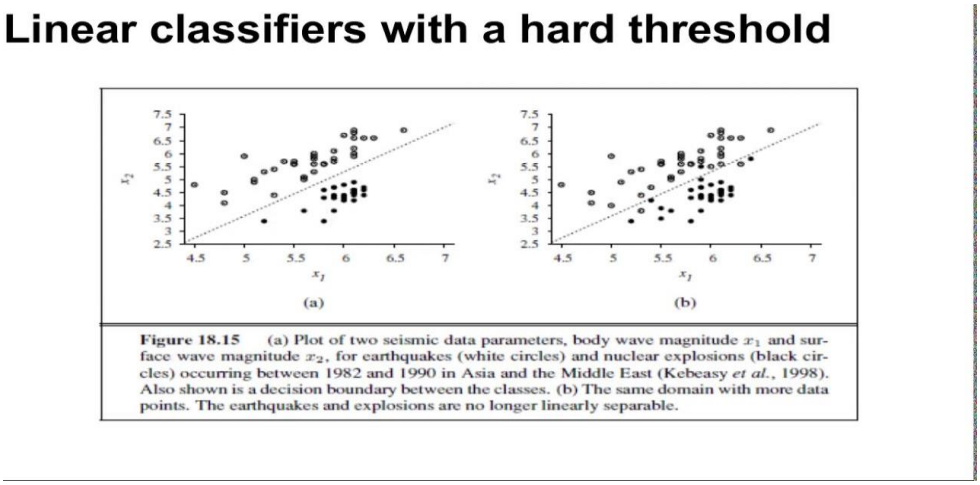


**Figure 18.15** (a) Plot of two seismic data parameters, body wave magnitude $x_1$ and surface wave magnitude $x_2$, for earthquakes (white circles) and nuclear explosions (black circles) occurring between 1982 and 1990 in Asia and the Middle East (Kebeasy *et al.*, 1998). Also shown is a decision boundary between the classes. (b) The same domain with more data points. The earthquakes and explosions are no longer linearly separable.

Mathematically, the model computes:

z=w0+w1X1+w2X2+⋯+wnXnz = w_0 + w_1X_1 + w_2X_2 + \dots + w_nX_nz=w0+w1X1+w2X2+⋯+wnXn

where wiw_iwi are the weights, XiX_iXi are the input features, and w0w_0w0 is the bias or intercept.

Then, a **hard threshold** is applied, such as:

output={1if z≥00if z<0\text{output} = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases}output={10if z≥0if z<0

This means the model predicts class **1** if the weighted sum is zero or positive, otherwise class **0**.

### Example:

Suppose you want to classify emails as spam or not spam based on two features:

- X1X_1X1: number of suspicious words
- X2X_2X2: number of links

The linear classifier might compute:

z=−1+0.5×X1+0.3×X2z = -1 + 0.5 \times X_1 + 0.3 \times X_2z=−1+0.5×X1+0.3×X2

If z≥0z \geq 0z≥0, classify the email as spam (1); otherwise, not spam (0).

## Linear classification with logistic Regression:

**Linear classification with logistic regression** is a widely used technique for binary classification problems, where the goal is to assign inputs to one of two classes (e.g., spam vs. not spam). Unlike a linear classifier with a hard threshold that outputs a strict 0 or 1 based on whether the weighted sum of inputs crosses zero, logistic regression models the **probability** that an input belongs to a particular class.

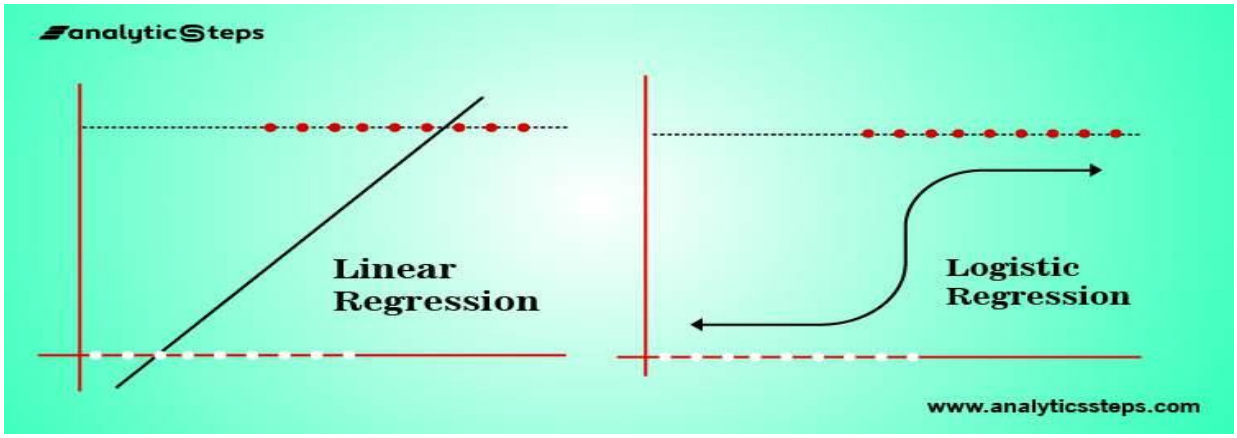


The core idea is to first compute a linear combination of the input features:

z=w0+w1X1+w2X2+⋯+wnXnz = w_0 + w_1X_1 + w_2X_2 + \dots + w_nX_nz=w0+w1X1+w2X2+⋯+wnXn

Then, instead of applying a hard cutoff, logistic regression passes this value through the **sigmoid function**:

σ(z)=11+e−z\sigma(z) = \frac{1}{1 + e^{-z}}σ(z)=1+e−z1

The sigmoid function transforms any real number into a value between 0 and 1, which can be interpreted as the **probability** that the input belongs to class 1. If the probability is above a chosen threshold (commonly 0.5), the input is classified as class 1; otherwise, class 0.

**Example:**

Imagine you want to classify whether a tumor is malignant or benign based on features like size and texture. Logistic regression computes a weighted sum of these features and then applies the sigmoid function to estimate the probability the tumor is malignant.

If the model outputs a probability of 0.8 for a tumor, it means there's an 80% chance the tumor is malignant, and it would be classified as such if using the 0.5 threshold.

### Advantages of Logistic Regression:

- Produces probabilistic outputs, offering confidence levels for predictions.
- Can handle linearly separable data effectively.
- Optimized using maximum likelihood estimation, providing well-calibrated models.
- Extensible to multiclass problems with variants like multinomial logistic regression.

## ARTIFICIAL NEURAL NETWORKS:

An **Artificial Neural Network (ANN)** is a computational model inspired by the way biological neural networks in the human brain process information. It consists of layers of interconnected processing units called neurons, which work together to recognize patterns and solve problems. ANNs are widely used in artificial intelligence and machine learning tasks such as image recognition, speech processing, and language translation. They are particularly effective when dealing with large amounts of data and complex relationships.

The structure of an ANN is made up of three main types of layers: the **input layer**, one or more **hidden layers**, and an **output layer**. The input layer receives the raw data (such as pixel values from an image), and each neuron in this layer represents one feature of the data. The data is then passed through the hidden layers, where the network processes and transforms the information. These layers are called "hidden" because they are not directly exposed to the input or output. Each connection between neurons has a **weight** and a **bias**, which determine the importance of input values and how the neuron activates.

As the data moves through the network in a process called **forward propagation**, each neuron calculates a weighted sum of its inputs, adds a bias, and applies an **activation function**. This function introduces non-linearity into the model, allowing it to learn more complex patterns. Common activation functions include the sigmoid, tanh, and ReLU (Rectified Linear Unit). After processing through the network, the output layer provides the final result, such as a classification label or a numerical prediction.

To train an ANN, a technique called **backpropagation** is used. After the network makes a prediction, a **loss function** measures how far off the prediction is from the actual value. The network then adjusts its weights and biases to minimize this loss, using an optimization method like **gradient descent**. This process repeats many times, gradually improving the network's accuracy.

Artificial Neural Networks are powerful tools capable of learning from data, identifying patterns, and making intelligent decisions. They form the backbone of many modern AI applications and are a key part of deep learning, where networks have many hidden layers to capture high-level features in data.

**EXAMPLE :**

### ? Example: Predicting Whether a Student Will Pass an Exam

Imagine we want to predict whether a student will pass an exam based on two inputs:

- **Hours Studied**
- **Sleep Hours**

We use an ANN with:

- **2 input neurons** (for the two features),
- **1 hidden layer** with **2 neurons**,
- **1 output neuron** (which gives a value between 0 and 1, where >0.5 means "Pass").

## Diagram of the ANN

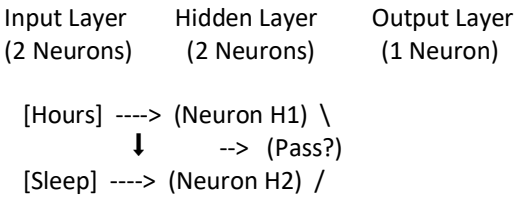


## How It Works Step-by-Step

1. **Inputs**:
   - Hours Studied = 5
   - Sleep Hours = 7
2. **Forward Propagation**:
   - Each hidden neuron (H1 and H2) receives a weighted sum of the inputs plus a bias.
   - Activation functions are applied (e.g., ReLU or sigmoid).
   - The output neuron processes the outputs of H1 and H2, applies another activation function, and gives a final prediction.
3. **Output**:
   - Let's say the final output is **0.82**, which we interpret as an **82% chance of passing**.
   - Since 0.82 > 0.5, we predict: **Pass**.

SINGLE LAYER AND MULTILAYER FORWARD NETWORKS

## SINGLE LAYER AND MULTILAYER FORWARD NETWORKS

**Single Layer Forward Networks**, often called **Single Layer Perceptrons**, are the simplest type of artificial neural networks. They consist of a single layer of neurons (or units) that take input features, apply weights, sum them up, and pass the result through an activation function, usually a hard threshold or a sigmoid for classification tasks. These networks can only solve problems that are **linearly separable**, meaning the data can be separated by a single straight line or hyperplane. For example, a single-layer perceptron can classify whether an email is spam or not based on simple linear combinations of features like the number of suspicious words or links. However, they fail on problems like the XOR problem where the classes are not linearly separable.

**Multilayer Forward Networks**, commonly known as **Multilayer Perceptrons (MLPs)** or **feedforward neural networks**, consist of multiple layers of neurons: an input layer, one or more hidden layers, and an output layer. Each layer is fully connected to the next, meaning every neuron in one layer sends its output to every neuron in the following layer. The main advantage of multilayer networks is their ability to model **complex, non-linear relationships** by stacking layers and applying non-linear activation functions (like ReLU, sigmoid, or tanh). The hidden layers enable the network to learn features at different levels of abstraction. For example, in image recognition, the first hidden layer might learn edges, the next layer might learn shapes, and deeper layers might recognize objects. Multilayer networks are trained using algorithms like **backpropagation** and **gradient descent** to adjust weights for accurate predictions. These networks form the backbone of modern deep learning.

## SUPPORT VECTOR MACHINES

**Support Vector Machines (SVMs)** are supervised machine learning algorithms used primarily for classification tasks, although they can also handle regression and outlier detection. The main idea behind SVM is to find the optimal boundary—called a **hyperplane**—that best separates the data into different classes. In a two-dimensional space, this hyperplane is simply a line, while in higher-dimensional spaces, it becomes a plane or a more complex structure. SVM aims to maximize the **margin**, which is the distance between the hyperplane and the nearest data points from each class. These closest data points are known as **support vectors**, and they are critical because they directly influence the position and orientation of the hyperplane. By focusing only on these key points, SVM achieves better generalization and avoids overfitting.

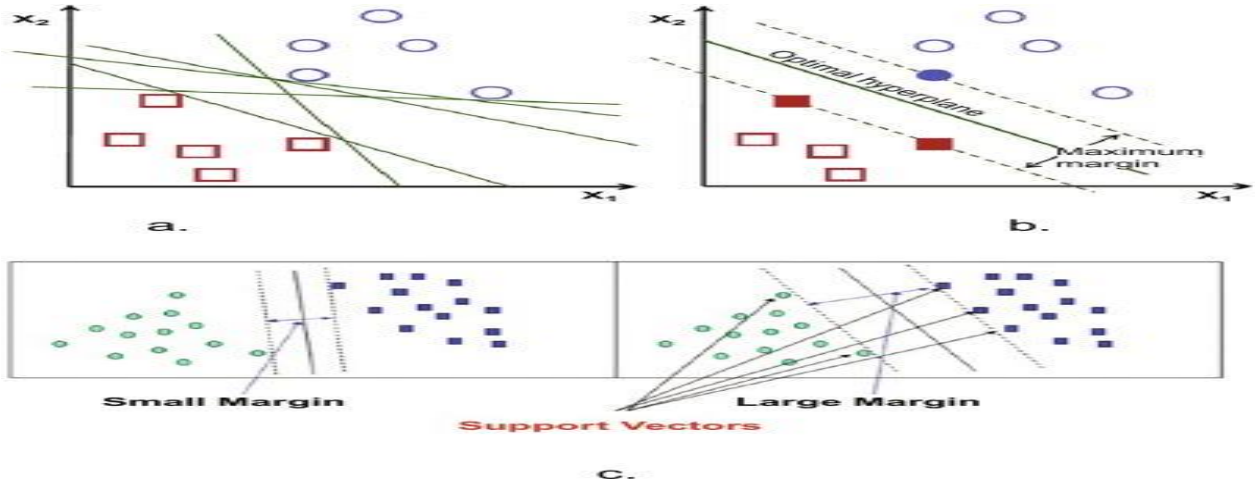


When data is linearly separable, a **linear SVM** works well by drawing a straight line (or hyperplane) to separate classes. However, in many real-world problems, the data is not linearly separable. To solve this, SVM uses a technique known as the **kernel trick**. Kernels transform the original data into a higher-dimensional space where a linear separation is possible. Common kernel functions include the **polynomial kernel**, **radial basis function (RBF)**, and **sigmoid kernel**. For example, imagine trying to separate points arranged in concentric circles. In two dimensions, no straight line can separate them. But by applying a kernel that projects the data into a third dimension, SVM can find a plane that separates the classes effectively.

A practical example of SVM is in **email spam detection**. Suppose we have a dataset of emails, each labeled as either "spam" or "not spam." The emails are represented using features like word frequencies, presence of

specific keywords, etc. SVM can be trained on this data to learn the boundary that separates spam from non-spam emails. Once trained, the model can classify new emails accurately based on the learned hyperplane.

## Example: Email Spam Detection using SVM

Imagine you work for an email service provider and want to build a **spam filter** that can automatically classify emails as either **“spam”** or **“not spam.”** To do this, you collect a dataset of emails. Each email is represented by features such as the number of times certain keywords appear (e.g., "free", "win", "money"), the presence of links or attachments, and the email length. Each email in your dataset is labeled as either spam (1) or not spam (0). This becomes your **training data**.

Now, you train a **Support Vector Machine (SVM)** on this dataset. The SVM algorithm analyzes the features of the emails and finds the **optimal hyperplane** that separates spam emails from non-spam ones. In this case, the data might not be perfectly linearly separable—for example, some legitimate emails might use the word "free." So, instead of a straight line, the SVM uses a **kernel trick** (like the Radial Basis Function, or RBF kernel) to project the data into a higher-dimensional space. In this new space, the SVM finds a decision boundary that can separate the two categories more effectively.

Once trained, this SVM model can now classify new, unseen emails. Suppose a new email arrives with the words "win a free gift card" and has suspicious links. Based on the features extracted from this email, the SVM model compares it to the support vectors and determines on which side of the hyperplane the email lies. If it falls on the "spam" side, the email is filtered out automatically.

**********